

Ideally, I would like to have gathered way more data points. I was able to achieve an almost exact beta match at around $n = 300$ observations/samples. Below are my two attempts at the lasso regression algorithm implementation. The first uses essentially the exact same notation from the computer assignment pdf. However, it results in a lot of chatter. The second uses a separate implementation based on the beginning parts of the slides. In other words, I define the cost function similar to the gradient descent problem, and then add the lambda L1 regularization term. This graph came out much cleaner but does not seem to be affected by varying lambda.

