

The gradient descent algorithm I implemented successfully converged and predicted the correct weight and bias terms as shown by the plot and console log. The logic I used to implement the algorithm came directly from slides 9 and 10 and should be straightforward to understand.

```
# hint: please check the slides of the gradient descent (1-3, the 9th to 10th page)
#####
notes:
p=2
y size = n x 1 = [1000 x 1]
X size = n x (p+1) = [1000 x 2] ?
beta size = (p+1) x 1 = [2 x 1]

basically our output hb(x) = beta[0] + xbatch * beta[1]
"""

def gradient_descent(beta, lr, x_batch, y_batch):
    alpha = lr

    slope = beta[0]
    intercept = beta[1]

    hb = slope * x_batch + intercept
    cost = (1/2) * np.sum([val**2 for val in (hb - y_batch)]) # define output # define cost function aka J(beta)

    dslope = np.sum((hb - y_batch) * x_batch) # differentiate with respect to the weight component beta[0] aka slope
    dintercept = np.sum((hb - y_batch)) # differentiate with respect to the bias component beta[1] aka intercept

    slope_next = slope - alpha * dslope # put all the terms together to weight/slope for the next slope term
    intercept_next = intercept - alpha * dintercept # put all the terms together to bias/intercept for the next slope term

    weight_next = slope_next
    bias_next = intercept_next
    beta_next = [weight_next, bias_next]

    return cost, beta_next

#####
```

