

Modern Machine Learning

Computer Assignment #2

1. *Binary Logistic Regression:* Using MATLAB or Python, implement the cost function for logistic regression and then minimize it.

Background: The cost function to be minimized is the *negative* log-likelihood function

$$NLL(\boldsymbol{\beta}) = - \sum_{i=1}^n y_i \log(p(\mathbf{x}_i, \boldsymbol{\beta})) + (1 - y_i) \log(1 - p(\mathbf{x}_i, \boldsymbol{\beta})),$$

where $p(\mathbf{x}_i, \boldsymbol{\beta}) = \frac{1}{1+e^{-\boldsymbol{\beta}^T \mathbf{x}_i}}$ (logistic function). The derivative of the negative log-likelihood can be expressed as

$$\frac{\partial NLL(\boldsymbol{\beta})}{\partial \boldsymbol{\beta}} = - \sum_{i=1}^n \mathbf{x}_i (y_i - p(\mathbf{x}_i, \boldsymbol{\beta}))$$

Python files: The `log_regression_example.py` uses the Tensorflow package to realize the logistic regression for binary classification based on randomly generated data. It also plots the data and the decision boundary.

Submission guidelines: Your submission should include:

- A unique **zip folder**, which should include a modified version of `log_regression.py`, in which you need to realize the function *logistic_regression*. The function *logistic_regression* is designed to minimizing the cost function $NLL(\boldsymbol{\beta})$ via gradient descent.
- the structure of *logistic_regression(beta, lr, x_batch, y_batch)* is summarized as follows: (i) **beta** is a 3×1 vector, (ii) *lr* denotes the learning rate, (iii) **x_batch** is the dataset of two linear models, the dimension of **x_batch** is 2000×3 , which means there are 2000 data points $\{(x, y, 1)\}_{i=1}^{2000}$, we extend $\{(x, y)\}_{i=1}^{2000}$ to $\{(x, y, 1)\}_{i=1}^{2000}$ for implement the matrix operation **x_batch** · **beta**, where beta_3 is the bias, e.g., $x \cdot \text{beta}_1 + y \cdot \text{beta}_2 + \text{beta}_3 \cdot 1$. In addition, **y_batch** is the training labels for **x_batch**. If (x, y) belongs to the first linear model, then $y_{batch} = 1$, otherwise $y_{batch} = 0$.

- Please rename the modified file `lin_regression.py` by adding your last name to the script name, e.g., `lin_regression_smith.py`.
- A pdf file with a comparative figure of the two lines obtained by your function and the `log_regression_example.py`. Explain possible differences.

MATLAB files: The script `log_regression_example.m` uses the MATLAB function `glmfit(·)` to obtain the logistic regression coefficients for binary classification using randomly generated data. It also plots the data and the decision boundary.

Submission guidelines: Your submission should include:

- A unique **zip folder**, which should include a modified version of `log_regression_example.m` and `costLogistic.m`. The structure of the function `costLogistic.m` is the following: $[NLL(\boldsymbol{\beta}), \frac{\partial NLL(\boldsymbol{\beta})}{\partial \boldsymbol{\beta}}] = \text{costLogistic}(\boldsymbol{\beta}, \mathbf{X}, \mathbf{y})$, where \mathbf{X} are the observations, $\boldsymbol{\beta}$ are the logistic regression coefficients, and \mathbf{y} are the class labels of the examples in \mathbf{X} . Make sure the results you obtain are very close to the ones returned by the function `glmfit()`. Please rename the modified file `log_regression_example`, replacing the word 'example' in the provided script with your last name. For example `log_regression_smith.m`. **This should be the main function.**
 - A pdf file with a comparative figure of the two lines obtained by your function and the function `glmfit()`. Explain possible differences.
2. *Multiclass classification and regularization:* Using MATLAB, implement the *one-vs-all* technique for multiclass classification using regularized logistic regression.

Background: The one-vs-all approach consists of training K separate binary classifiers, where K is the number of classes, for each one of the different classes.

- Train K separate binary classifiers producing $\boldsymbol{\beta}_j$ ($j = 1, 2, \dots, K$)
- For every new example \mathbf{x} , find the predicted label as

$$\hat{j} = \arg \max_j p(\mathbf{x}, \boldsymbol{\beta}_j), \text{ where } p(\mathbf{x}_i, \boldsymbol{\beta}) = \frac{1}{1 + e^{-\boldsymbol{\beta}^T \mathbf{x}_i}}$$

Recall that the equations for regularized logistic regression are the following

$$NLL(\boldsymbol{\beta}) = - \sum_{i=1}^n y_i \log(p(\mathbf{x}_i, \boldsymbol{\beta})) + (1 - y_i) \log(1 - p(\mathbf{x}_i, \boldsymbol{\beta})) + \frac{\lambda}{2} \sum_{j=1}^m \beta_j^2,$$



Figure 1: Example images from MNIST database

where m is the number of variables, and $p(\mathbf{x}_i, \boldsymbol{\beta})$ is the logistic function defined above.

$$\frac{\partial NLL(\boldsymbol{\beta})}{\partial \beta_j} = - \sum_{i=1}^n x_i^{(j)} (y_i - p(\mathbf{x}_i, \boldsymbol{\beta})) + \lambda \beta_j$$

where $x_i^{(j)}$ is the j -th component of vector \mathbf{x}_i .

The data used in this exercise is a portion of the MNIST database. It contains handwritten digits 0-9. The size of the images is 28×28 pixels, these pixels are vectorized producing features of size 784×1 . Your MATLAB program will attempt to recognize these digits. Examples of the handwritten digits can be seen in Figure 1.

Python files: The `multiclass_log_reg.py` is the template to implement your one-vs-all classifier.

Submission guidelines: Your submission should include:

- To run the `multiclass_log_reg.py` script, you are required to install **scipy** package into your Anaconda environment. The command is **conda install scipy**.
- A unique **zip folder**, which should include a modified version of `multiclass_log_reg.py`, in which you need to realize the function *logistic_regression*. The function *logistic_regression* is designed to minimizing the cost function $NLL(\boldsymbol{\beta})$ with $L2$ norm regularization.
- the structure of *logistic_regression(beta, lr, x_batch, y_batch, lambda)* is summarized as follows: (i) **beta** is a 785×1 vector, (ii) *lr* is the learning rate, (iii) **x_batch** is the MNIST dataset consisting 5000 images, the dimension of **x_batch** is 5000×785 , which means there are 5000 digit images $\{(\mathbf{x}, 1)\}_{i=1}^{5000}$, we extend $\{\mathbf{x}\}_{i=1}^{5000}$ to $\{(\mathbf{x}, 1)\}_{i=1}^{5000}$ for implement the matrix operation $\mathbf{x_batch} \cdot \mathbf{beta}$, where \mathbf{beta}_{785} is the bias, e.g., $\sum_{i=1}^{784} x_i \cdot \mathbf{beta}_i + \mathbf{beta}_{785} \cdot 1$. **y_batch** is the training labels for **x_batch**, its dimension is 5000×10 . If \mathbf{x} belongs to the i th digit, then $y_{batch} = [0, \dots, 1, \dots, 0]$.
- Please rename the modified file `logistic_regression.py` by adding your last name to the script name, e.g., `lin_regression_smith.py`.

- A pdf file with a figure showing the classification accuracy vs the regularization parameter λ . Vary λ from 0 to 20, in steps of 1. Explain your results.

MATLAB files: The script `multiclass_log_reg_example.m` can be used as a template to implement your one-vs-all classifier. **The function `fmincg.m` is also provided to speed up the optimization process.** This function works the same way as the function `fminunc()` provided by MATLAB, but it is faster.

Submission guidelines: Your submission should include:

- A unique **zip folder**, which should include a modified version of `multiclass_log_reg_example.m` and `multiclassLog.m`. The structure of the function `multiclassLog.m` is the following: $\hat{\beta}_{matrix} = \text{multiclassLog}(\mathbf{X}, \mathbf{y}, K)$, where \mathbf{X} are the observations, $\hat{\beta}_{matrix}$ are the K classifiers stacked column-wise as a matrix, and K is the number of class labels, in this case $K = 10$. This function should train K binary classifiers using your past implementation of logistic regression. You can use a *for* loop to achieve this. Be careful setting the labels so that for class j , the training examples belonging to class j will be assigned label 1, and the other examples label 0. The function `pred = predictmulticlass($\hat{\beta}_{matrix}$, \mathbf{y})` is provided to predict the class labels using $\hat{\beta}_{matrix}$.

Please rename the modified file `multiclass_log_reg_example.m` replacing the word 'example' in the provided script with your last name. As in the previous exercise. **This should be the main function.**

- A pdf file with a figure showing the classification accuracy vs the regularization parameter λ . Vary λ from 0 to 20, in steps of 1. Explain your results.

Hint: For $\lambda = 0$, the classification accuracy should be about 87.5%.